

# Transfer Learning

## Basic concepts and code examples



**MAC 2014-2020**  
Cooperación Territorial

**Interreg**  
Fondo Europeo de Desarrollo Regional



EUROPEAN UNION



UNIVERSIDAD DE LAS PALMAS  
DE GRAN CANARIA



Tecnología Médica  
para el Desarrollo Sostenible

**1**

**Basic concepts**

# 1.1 Transfer learning: Definition




## Reminder

In section 3 we have mentioned that the traditional training process consists of **initializing the network parameters randomly** and optimizing them based on the dataset.



However, this process can be **inefficient** when a sufficiently large dataset is not available to train the model on a problem.

Transfer learning is a technique in which we use **models that have been previously trained** (usually with large databases) to **apply them to a specific dataset of interest**. Working with pre-trained networks allows us to take advantage of the knowledge acquired by these models and apply it to another related task.



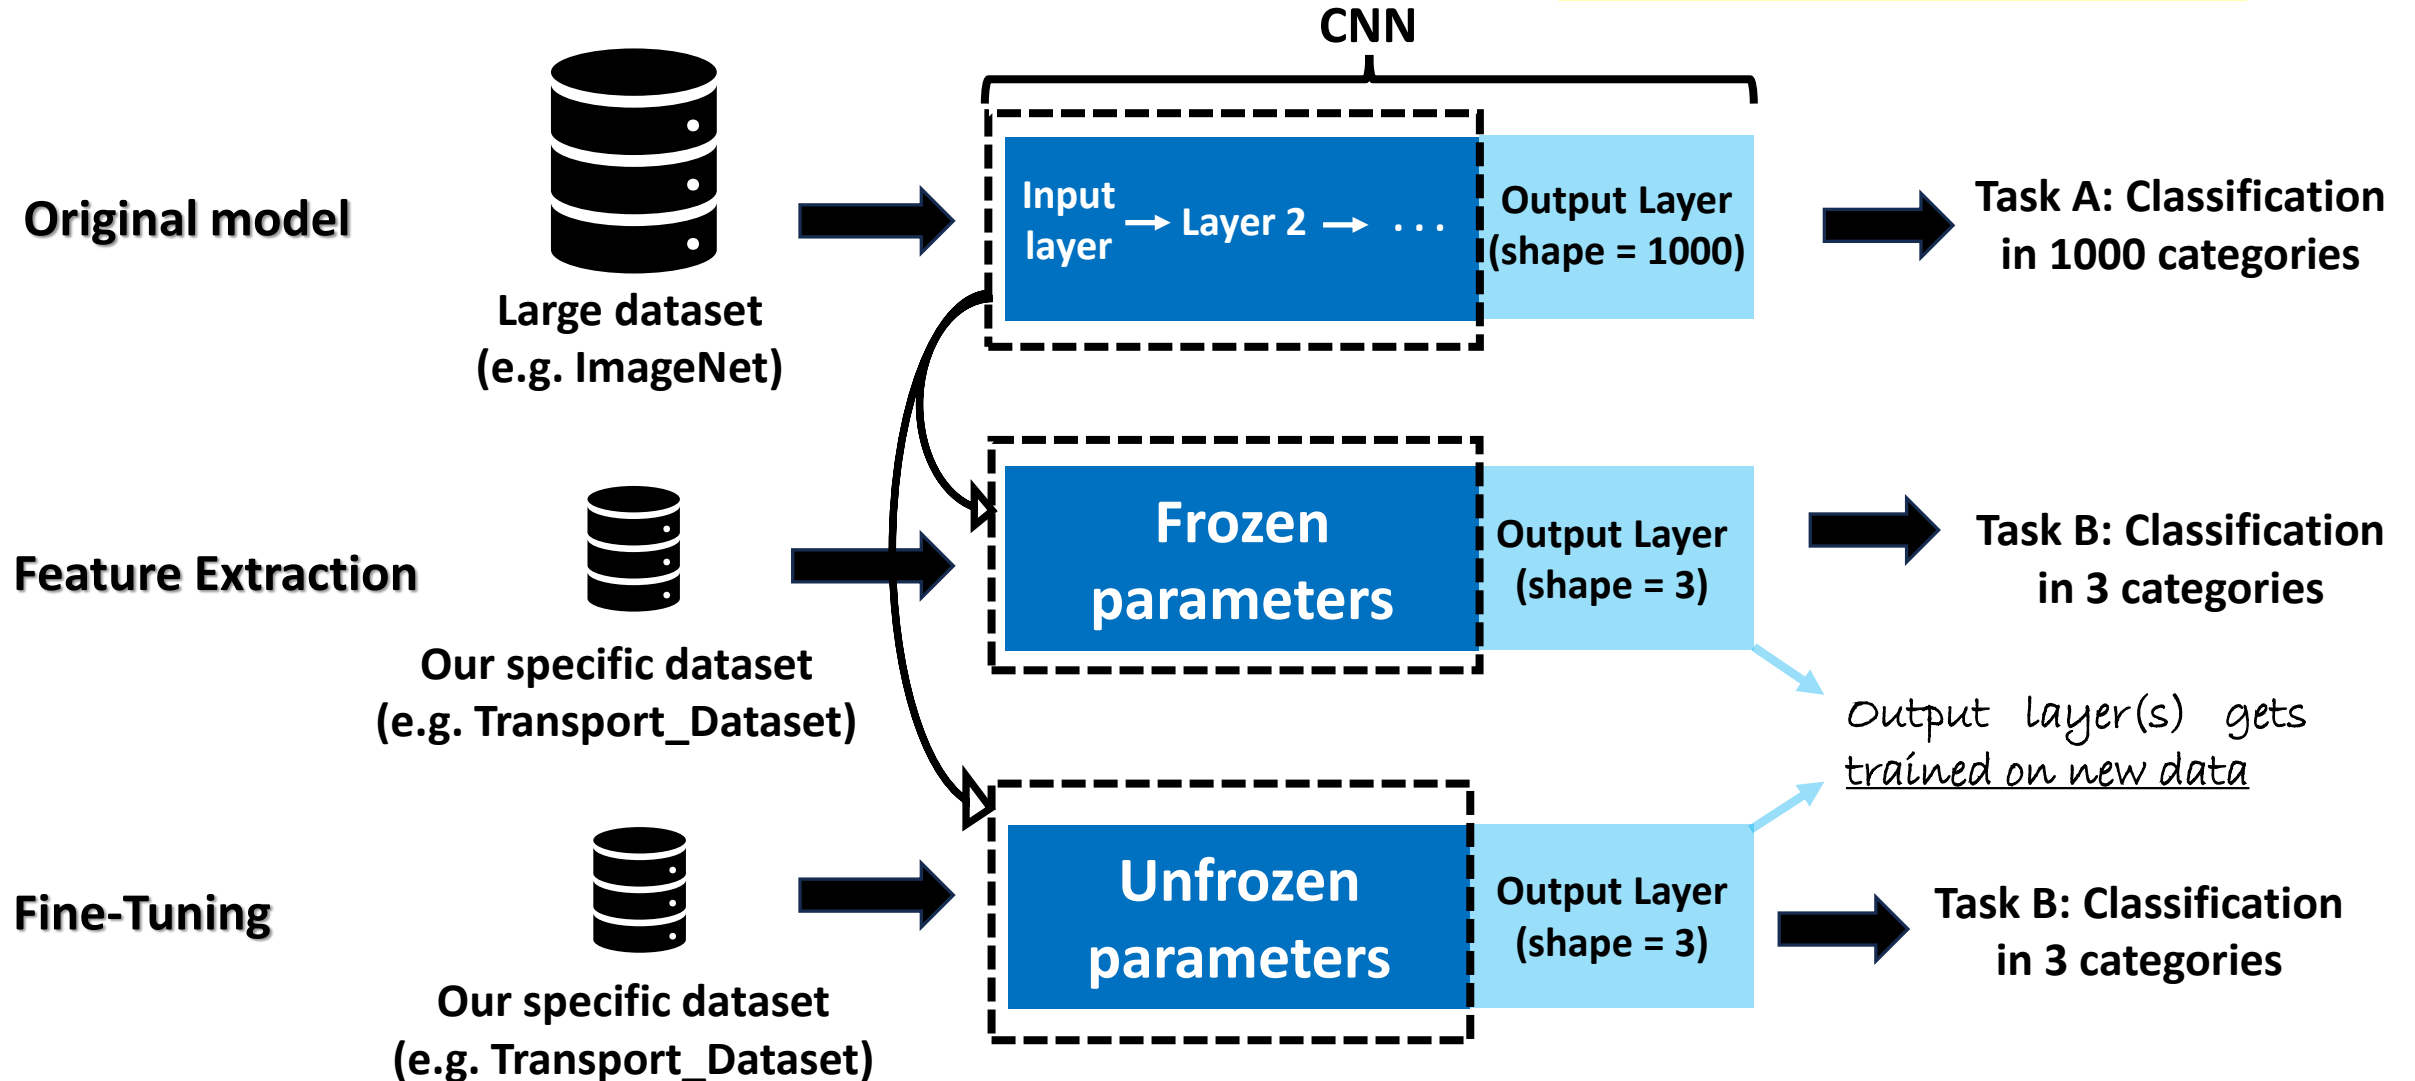
When **transfer learning** is used, the model **performs better** than when **trained from scratch** (with random initialization)

# 1.2 Feature extraction and fine-tuning

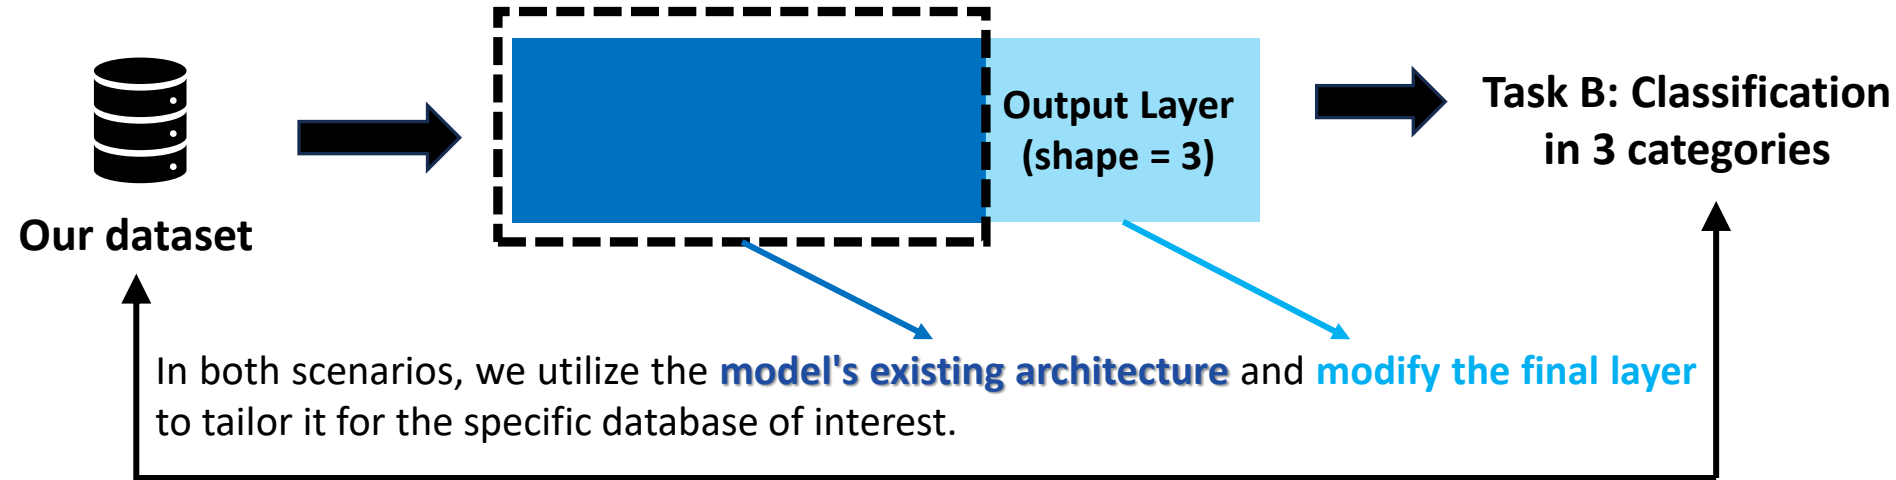


There are mainly two approaches to transfer knowledge from one model to another: **Feature extraction** and **Fine tuning**.

**“Frozen”**: original model layers don't update during training



# 1.2 Feature extraction and fine-tuning



**Feature extraction:** During feature extraction, the weights of the initial layers are frozen, meaning they are not updated during the training process on the new task. Only the additional layers (usually a new classifier) added on top of the pre-trained model are trained on the new data.

**Fine-tuning:** It involves updating the weights of many or all the layers of the pre-trained model, including both the feature extractor and the classifier layers, on the new task-specific data. Fine-tuning usually requires more data than feature extraction

**2**


# **Transfer learning in PyTorch**

## 2.1 Get a model



There are several places you can find pretrained models to use for your own problems:

- PyTorch domains libraries ( e.g [torchvision.models](#))
- [HuggingFace Hub](#)
- [Torch Image Models \(timm library\)](#)



```
import torchvision
model = torchvision.models.resnet18(weights = 'IMAGENET1K_V1')
print(model)
```

The original model is pretrained on a large dataset, such as ImageNet

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    ...
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Linear(in_features=512, out_features=1000, bias=True)
  )
)
```

Extracts features from image.  
We will maintain the model's existing architecture

Turns features into a feature vector

Utilizing the feature vector, the model generates predictions for 1000 classes

## 2.2 Replace the Fully Connected layer



Substitute the last layer of the model for a classification into three categories:

```
num_ftrs = model.fc.in_features
model.fc = nn.Linear(num_ftrs, 3)
```

```
(avgpool): AdaptiveAvgPool2d(output size=(1, 1))
(fc): Linear(in_features=512, out_features=3, bias=True)
```

Output Layer  
(shape = 3)

To learn more about our model, let's use [torchinfo.summary\(\)](#)

Layer (type (var_name))	Input Shape	Output Shape	Param #	Trainable
ResNet (ResNet)	[32, 3, 224, 224]	[32, 2]	--	True
└Conv2d (conv1)	[32, 3, 224, 224]	[32, 64, 112, 112]	9,408	True
└BatchNorm2d (bn1)	[32, 64, 112, 112]	[32, 64, 112, 112]	128	True
└ReLU (relu)	[32, 64, 112, 112]	[32, 64, 112, 112]	--	--
└MaxPool2d (maxpool)	[32, 64, 112, 112]	[32, 64, 56, 56]	--	--
└Sequential (layer1)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	True
└BasicBlock (0)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	True
└Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	36,864	True
└BatchNorm2d (bn1)	[32, 64, 56, 56]	[32, 64, 56, 56]	128	True
...				
└AdaptiveAvgPool2d (avgpool)	[32, 512, 7, 7]	[32, 512, 1, 1]	--	--
└Linear (fc)	[32, 512]	[32, 2]	1,026	True

Input and output shape  
of data per layer

Informs if the layers  
are unfrozen

```
Total params: 11,177,538
Trainable params: 11,177,538
Non-trainable params: 0
Total mult-adds (G): 58.03
```

Total number of parameters and  
trainable parameters

The entire model is trainable:  
Fine tuning



## 2.3 Choose the trainable layers/parameters



```
model = torchvision.models.resnet18(weights = 'IMAGENET1K_V1')
```

```
for param in model.parameters():  
    param.requires_grad = False
```

Freezing the model's parameters to avoid updating them during training

```
num_ftrs = model.fc.in_features  
model.fc = nn.Linear(num_ftrs, 3)
```

Layer (type (var_name))	Input Shape	Output Shape	Param #	Trainable
ResNet (ResNet)	[32, 3, 224, 224]	[32, 2]	--	Partial
└─Conv2d (conv1)	[32, 3, 224, 224]	[32, 64, 112, 112]	(9,408)	False
└─BatchNorm2d (bn1)	[32, 64, 112, 112]	[32, 64, 112, 112]	(128)	False
└─ReLU (relu)	[32, 64, 112, 112]	[32, 64, 112, 112]	--	--
└─MaxPool2d (maxpool)	[32, 64, 112, 112]	[32, 64, 56, 56]	--	--
└─Sequential (layer1)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─BasicBlock (0)	[32, 64, 56, 56]	[32, 64, 56, 56]	--	False
└─Conv2d (conv1)	[32, 64, 56, 56]	[32, 64, 56, 56]	(36,864)	False
...				
└─AdaptiveAvgPool2d (avgpool)	[32, 512, 7, 7]	[32, 512, 1, 1]	--	--
└─Linear (fc)	[32, 512]	[32, 3]	1,539	True

Most layers are untrainable (frozen)

```
Total params: 11,178,051  
Trainable params: 1,539  
Non-trainable params: 11,176,512  
Total mult-adds (G): 58.03
```

Less trainable parameters (That is why feature extraction tends to result in faster training)

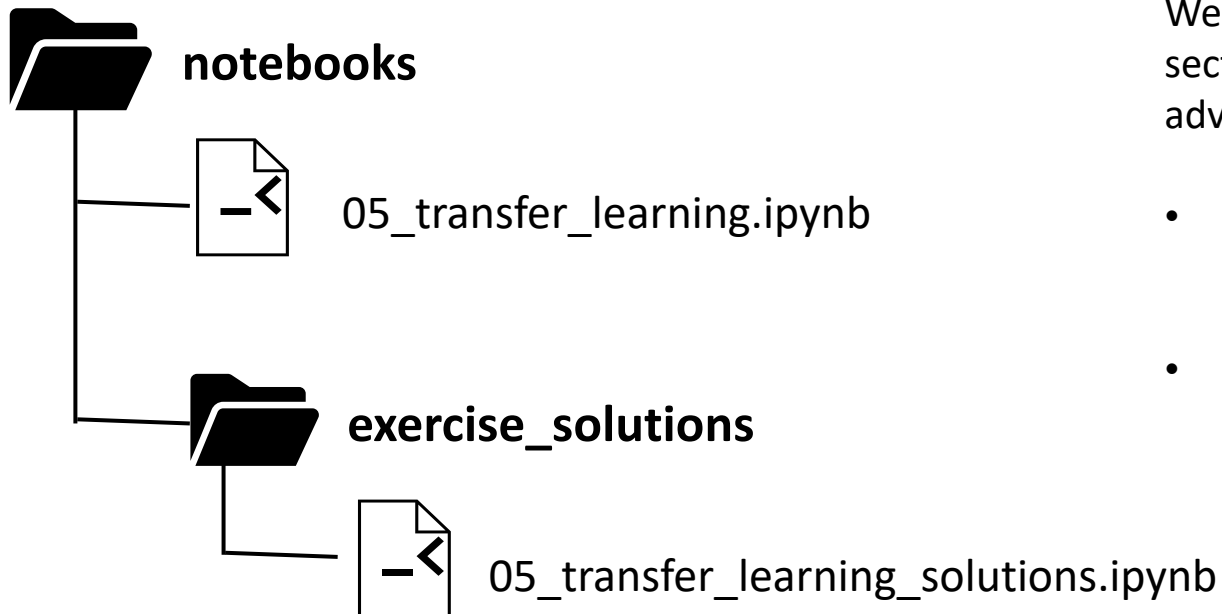
Only last layer is trainable: Feature extraction

## 2.2 Code examples



### Let's code!

To gain hands-on experience with transfer learning, we will utilize a pretrained ResNet-18 model on the Imagenet database and fine-tune it for our specific 4-class classification problem. Subsequently, we will examine some inference examples and assess the performance of the model.



We will see that, compared to the training carried out in section 3, the use of Transfer learning has 2 fundamental advantages:

- It can leverage an **existing neural network architecture** proven to work on problems similar to the current task.
- It can leverage a working network architecture that **has already learned patterns from similar data** to the current task, often resulting in **superior results with less data**.